
Reasoning Depth Erodes the Premise Verification That Post Training Builds

Anonymous Author(s)

Affiliation

Address

email

Abstract

Whether a language model can tell that a problem does not determine an answer has been measured on released models, which cannot say when the ability arrives in training or what takes it away. Instella publishes intermediate checkpoints spanning pretraining and post training together with its full data mixture, so both questions become measurable. A premise deletion probe removes exactly one premise the reference solution consumes, giving 3179 variants over 1779 parent problems, and every abstention rate is read against an answerable control on the untouched parents, where declining is an error. Verification is close to absent before post training and is built by it: discrimination runs +1.84, +10.10, +38.08 and +52.63 points across stage one, stage two, supervised fine tuning and instruct, while wrong refusals stay below 3.5 percent throughout, and a reworded licence preserves the ladder. An independently built MATH probe reproduces it, and three instruct checkpoints from other families all discriminate. Reasoning depth then erodes what training builds. Abstention on underdetermined problems falls 5.22 points per additional calculator step while the answerable control stays flat at -0.11, a difference of -5.11 on $[-6.36, -3.82]$, and the decay steepens as the checkpoint improves. Under the answer forcing prompt used for benchmark scoring the strongest checkpoint still answers 97.6 percent of provably underdetermined items. Controlled injection at a fixed token budget drives verbatim reproduction to 0.946 and returns the memorised calculation chain, at +5.44 on $[+1.46, +9.38]$, without returning the answer, at +1.52 on $[-2.12, +5.21]$, so what depth erodes is not recall of training text.

1 Introduction

A model that answers a grade school arithmetic problem has been given premises that determine an answer. Whether it checked that they do is invisible to the score, because every item on a benchmark is answerable. Probing that directly requires items that are not, and a growing body of work now supplies them. QuestBench casts underspecified reasoning tasks as constraint satisfaction problems with a missing variable [Li et al., 2025a], AbstentionBench assembles unanswerable questions across domains and reports that reasoning fine tuning degrades abstention [Kirichenko et al., 2025], PCBench evaluates premise critique across fifteen models and finds that most detect a flawed premise only when explicitly prompted to look [Li et al., 2025b], and ReliableMath does the same for unsolvable mathematics [Xue et al., 2025]. That models often fail to verify their premises, and do better when told they may decline, is therefore established rather than at issue here.

What that work cannot ask is developmental. Every model it evaluates is a finished artifact, so the capability can be measured but not located: nothing says whether verification arrives with scale, with instruction tuning, or with preference optimisation, because the intermediate checkpoints of the models under test are not released. Instella [Liu et al., 2025] publishes four checkpoints spanning its pipeline along with the complete data mixture, which makes the question answerable on one family. This paper asks when premise verification appears, and what it is lost to.

Table 1: Declining with its control, under a prompt that names abstention as an option. Pruned items are underdetermined so declining is correct; parent problems are answerable so declining is an error. The last two columns give balanced accuracy of the decline decision, where 0.500 is chance, measured on the GSM8K probe and on an independently built MATH probe of 818 parents. Cluster bootstrap over parent problems, 4000 draws, seed 6198.

Checkpoint	GSM8K probe				Decision quality	
	Declines, und.	Declines, ans.	Discrimination	95% CI	GSM8K	MATH
Stage 1	0.043	0.025	+1.84	[+0.97, +2.70]	0.509	0.516
Instella 3B	0.110	0.008	+10.10	[+8.94, +11.30]	0.550	0.557
SFT	0.416	0.035	+38.08	[+36.06, +40.02]	0.690	0.611
Instruct	0.538	0.011	+ 52.63	[+50.77, +54.57]	0.763	0.675

Both answers turn out to concern reasoning rather than compliance. Verification is close to absent in the pretrained checkpoints and is manufactured by post training, rising from +1.84 to +52.63 points of discrimination against an answerable control. It then erodes with the depth of the chain the model is tracking, falling 5.22 points per additional reasoning step while the control stays flat, and eroding faster in the checkpoints that verify best. Because the corpus is published, the erosion can also be separated from memorisation: continued pretraining that drives verbatim reproduction of injected documents to 0.946 returns the memorised calculation chain without returning the answer, so what fails at depth is the check, not the recall.

2 Setup

Exposure. Containment is the fraction of an item’s thirteen grams appearing in a single corpus document, taking the maximum over documents. Measured on identical items through one extraction path, 9.65 percent of GSM8K train parents reach containment 0.999 or above against the 119,014 row subset consumed during stage two and 18.52 percent against the full pool, so the pool overstates exposure and the subset is the correct reference.

The probe. Each item removes one premise the reference solution consumes, verified by requiring the annotated calculator chain to depend on it. The result is 3179 variants over 1779 parent problems, split 1591 train and 1588 test. Every item is underdetermined by construction, so declining is the correct response and any specific number is wrong.

Checkpoints. Four public Instella checkpoints span the pipeline: stage one, which precedes the synthetic corpus, stage two, supervised fine tuning and instruct. Decoding is greedy and held fixed; base checkpoints are prompted with four exemplars and no chat template, since a base model loops without them.

Inference. Intervals are cluster bootstraps over parent problems with 4000 draws at seed 6198, because several variants can descend from one problem and resampling rows would understate the interval.

3 Verification of premise sufficiency

3.1 Post training builds the capability

Under the answer forcing prompt used for benchmark scoring, the instruct checkpoint declines on 2.4 percent of items that provably do not determine an answer and commits to one on the remaining 97.6 percent, or 99.7 percent counting any emitted number; the three earlier checkpoints commit on 99.5 percent or more. Naming ##### unanswerable as an available response, and changing nothing else, raises declining to 53.8 percent at the instruct checkpoint. That rate alone is ambiguous, since every probe item is underdetermined and a model offered an escape hatch may simply take it. Running the identical prompt over the untouched parent problems, which are answerable and where declining is an error, gives 1.1 percent. The difference of +52.63 points on [+50.77, +54.57] is the quantity of interest: declining tracks whether the problem can actually be answered.




Figure 1: (a) balanced accuracy of the decline decision across the training pipeline, on GSM8K and on an independently built MATH probe, where 0.500 is chance. (b) decline rate against solution chain length at the instruct checkpoint, with the answerable control on the same axes; the control is flat, so the decay is specific to detecting that premises are missing. (c) the same decision quality for instruct checkpoints of three further families.

Table 1 shows the capability developing across the pipeline. Stage one, which precedes the synthetic corpus entirely, discriminates by less than two points. Each subsequent stage improves it, and wrong refusals on answerable problems never exceed 3.5 percent, so the gain does not come from a growing willingness to decline.

A second control rules out the remaining reading, that a pruned problem reads oddly and the model responds to perturbed text. Inserting an irrelevant twenty word sentence carrying an unused number leaves a problem fully answerable while perturbing its surface. Declining moves from 0.0113 clean to 0.0170 with an off topic insertion and 0.0125 with an in domain one, against 0.538 when a premise is actually removed. Rewording the licence while holding its information and its token fixed preserves the ladder, at +0.80, +6.64, +33.41 and +43.77 points with every interval excluding zero, and moves the level: Appendix E reports both.

3.2 Reasoning depth erodes it

The capability is not uniform over problems. Figure 1 plots abstention against the length of the reference solution’s calculator chain. On underdetermined problems at the instruct checkpoint the rate falls from 0.623 at one step to 0.364 at seven, a slope of -5.22 points per step. On the answerable control the slope is -0.11 with an interval containing zero, so the pattern is not a tendency to decline more on short problems. The difference of slopes is -5.11 on $[-6.36, -3.82]$.

The decay steepens as the checkpoint improves, running -2.34 , -3.64 and -5.11 for stage two, supervised fine tuning and instruct. Two alternatives are ruled out. Truncation never exceeds 0.7 percent at any chain length and restricting to terminated generations reproduces the rates exactly. Difficulty is separated by splitting parents on whether the model answers the intact problem correctly, which it does for 76.4 percent: the slope is -4.93 points per step on $[-6.51, -3.36]$ where it can solve the problem and -4.36 on $[-6.54, -2.19]$ where it cannot, so chain depth carries information difficulty does not account for.

3.3 Both results hold on a second domain and in other families

Two further sweeps address the single benchmark and single family. A premise deletion probe built independently on MATH [Hendrycks et al., 2021], 818 parents and 1091 variants drawn from all seven subjects under the same two validity checks, reproduces the ladder: decision quality runs 0.516, 0.557, 0.611 and 0.675 across the same checkpoints, against 0.509, 0.550, 0.690 and 0.763 on GSM8K. The shape is the same and the level is lower, which is expected, since competition problems are harder and a model declining one it simply cannot solve inflates the answerable rate and shrinks the measured gap. The depth decay reproduces there too, at -6.41 points per additional premise sentence on $[-8.43, -4.49]$, steepening across the pipeline as on GSM8K. The MATH control arm was built without a depth variable, so this replicates the gradient but not its control.

Across families, three further instruct checkpoints graded the same GSM8K items under the same prompt. All four discriminate, with decision quality of 0.763 for Instella 3B Instruct, 0.795 and

0.759 for Qwen2.5 at 3B and 1.5B, and 0.573 for OLMo-2 1B Instruct, so premise verification is a property of instruction tuned small language models rather than of one pipeline. The spread is not a ranking, since the models differ in pretraining corpus, instruction tuning recipe and training compute, and nothing in the design attributes a few points to any of those. OLMo-2 is the informative case: its discrimination of +14.68 points on [+12.13, +17.23] clears zero, yet it declines 29.2 percent of answerable problems, so its decisions sit near chance. A difference of rates would call that a weaker version of the same capability; the decision metric identifies it as declining freely rather than selectively, which is why false refusal is reported alongside discrimination throughout.

4 What exposure returns

4.1 Continued pretraining at controlled doses

Whether exposure supplies the answer is causal, so selected items were injected by continued pretraining at five doses under a fixed 8,388,608 token budget, held constant so a higher dose displaces filler rather than buying compute. Every arm carries positive controls evaluated before and after: injected document loss must fall and the model must continue an injected document verbatim, since a flat dose response with failed controls would measure a broken training run. Injection worked, verbatim reproduction climbing from 0.133 to 0.946 while held out loss rises, so the mixture taught those documents rather than the task.

4.2 The chain returns, the answer does not

Against a never injected control at the same dose, reproduction of the memorised calculation chain reaches +5.44 points on [+1.46, +9.38] at sixty four repetitions and rises with dose, while reproduction of the final answer reaches +1.52 on [-2.12, +5.21], an interval containing zero. The model reproduces the source text almost exactly and still does not recover the deleted quantity or the answer it determines. Repeating the identical mixture, budget, seed and probe on OLMo-2-1B, Qwen2.5-1.5B and Qwen2.5-3B gives +0.28, +0.00 and +2.28 points at the top dose, every interval containing zero, across ten dose points. Qwen2.5-3B is the sharpest case, driven to verbatim reproduction of 1.000 and still recovering the deleted answer no more often than a model that never saw the documents. Chain reproduction is scored on three of the four models, and is established on Instella and underpowered on the other two, as Appendix D sets out.

5 Discussion

Two capabilities that a single accuracy number conflates come apart here. Executing a solution procedure is present early and improves smoothly; verifying that the premises support a solution at all appears later and degrades with the depth of the chain being tracked. Benchmark scoring cannot see the difference, because every item it scores is answerable. Reporting declining under a prompt that permits it, against an answerable control, costs one evaluation pass and makes the second capability visible.

Limitations. One problem generator for the injection arm. Detection is measured through an explicit abstention option, so a model that recognises insufficiency but lacks the vocabulary to say so is scored as failing. Difficulty is stratified by the model’s own accuracy on intact problems rather than by an independent measure, and depth on MATH is sentence count rather than a verified calculation chain, which that benchmark does not annotate. Chain reproduction under injection is established on one model, and the replication is underpowered for it. Attribution to specific training documents is not established anywhere.

Conclusion. Measured against exposure fixed by containment rather than estimated, Instella’s grade school arithmetic performance is not recovered text: no test item approaches containment against the corpus, and heavy controlled injection returns the memorised calculation chain without returning the answer. What the probe reveals instead is a separable capability. The model detects that a problem no longer determines an answer, does so increasingly well across training, and loses that detection as the reasoning chain lengthens.

References

- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few shot learners. In *Advances in Neural Information Processing Systems*, 2020. arXiv:2005.14165.
- Renjie Gu, Jiaxu Li, Yihao Wang, Yun Yue, Hansong Xiao, Yefei Chen, Yuan Wang, Chunxiao Guo, Pei Wei, Jinjie Gu, and Yixin Cao. Bridging the detection to abstention gap in reasoning models under insufficient information. *arXiv preprint arXiv:2605.28070*, 2026.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2021.
- Kaixuan Huang, Jiacheng Guo, Zihao Li, Xiang Ji, Jiawei Ge, Wenzhe Li, Yingqing Guo, Tianle Cai, Hui Yuan, Runzhe Wang, Yue Wu, Ming Yin, Shange Tang, Yujia Huang, Chi Jin, Xiaotong Chen, Chi Zhang, and Mengdi Wang. MATH perturb: Benchmarking LLMs mathematical reasoning abilities against hard perturbations. In *International Conference on Machine Learning*, 2025. arXiv:2502.06453.
- Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. AbstentionBench: Reasoning LLMs fail on unanswerable questions. *arXiv preprint arXiv:2506.09038*, 2025.
- Belinda Z. Li, Been Kim, and Zi Wang. QuestBench: Can LLMs ask the right question to acquire information in reasoning tasks? In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2025a. arXiv:2503.22674.
- Jinze Li, Gengxu Li, Yi Chang, and Yuan Wu. Don’t take the premise for granted: Evaluating the premise critique ability of large language models. *arXiv preprint arXiv:2505.23715*, 2025b.
- Jiang Liu, Jialian Wu, Xiaodong Yu, Yusheng Su, Prakamya Mishra, Gowtham Ramesh, Sudhanshu Ranjan, Chaitanya Manem, Ximeng Sun, Ze Wang, Pratik Prabhanjan Brahma, Zicheng Liu, and Emad Barsoum. Instella: Fully open language models with stellar performance. *arXiv preprint arXiv:2511.10628*, 2025.
- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *International Conference on Machine Learning*, 2023. arXiv:2302.00093.
- Benedikt J. Wagner. Two axes of LLM abstention: Answer correctness and question answerability. *arXiv preprint arXiv:2607.08456*, 2026.
- Mingqi Wu, Zhihao Zhang, Qiang Dong, Zhiheng Xi, Jun Zhao, Senjie Jin, Xuanjing Fan, Yuan Zhou, Han Lv, Ming Zhang, Yuchen Fu, Qi Liu, Sheng Zhang, and Qi Zhang. Reasoning or memorization? unreliable results of reinforcement learning due to data contamination. *arXiv preprint arXiv:2507.10532*, 2025.
- Boyang Xue, Qi Zhu, Rui Wang, Sheng Wang, Hongru Wang, Minda Hu, Fei Mi, Yasheng Wang, Lifeng Shang, Qun Liu, and Kam-Fai Wong. ReliableMath: Benchmark of reliable mathematical reasoning on large language models. *arXiv preprint arXiv:2507.03133*, 2025.
- Hugh Zhang, Jeff Da, Dean Lee, Vaughn Robinson, Catherine Wu, Will Song, Tiffany Zhao, Pranav Raja, Dylan Slack, Qin Lyu, Sean Hendryx, Russell Kaplan, Michele Lunati, and Summer Yue. A careful examination of large language model performance on grade school arithmetic. *arXiv preprint arXiv:2405.00332*, 2024.

A Related work

Premise verification and abstention. That models answer questions which do not determine an answer is well documented. QuestBench formalises underspecified reasoning as a constraint satisfaction problem missing one variable and finds that models fail to identify the missing variable even where they solve the fully specified problem [Li et al., 2025a]. AbstentionBench gathers unanswerable questions across domains and reports that reasoning fine tuning degrades abstention rather than improving it [Kirichenko et al., 2025]. PCBench evaluates premise critique over fifteen models and finds that most detect a flawed premise only under explicit prompting, which is the same licensing dependence measured in Section 3.1 [Li et al., 2025b], and ReliableMath establishes the corresponding benchmark for unsolvable mathematics [Xue et al., 2025]. Concurrent work sharpens two pieces of this. Gu et al. [Gu et al., 2026] name the gap between detecting insufficiency and acting on it and train an explicit answerability decision to close it; Wagner [Wagner, 2026] separates confidence in an answer from confidence that a question is answerable and shows the usual confidence signal is nearly blind to the second. The present work differs in what it varies. Those studies evaluate released models, so they can establish that the capability is uneven but not where it comes from or what removes it. Measuring four checkpoints of one published pipeline turns the capability into a developmental quantity, and reports two gradients rather than a level: it is manufactured almost entirely by post training, and it decays with the depth of the reasoning chain, faster in the checkpoints that verify best. The AbstentionBench finding that reasoning fine tuning degrades abstention and the ladder reported here are not in conflict on this reading, since the depth gradient is a within checkpoint effect that a single benchmark average over problems of mixed depth would partially absorb.

Memorisation and contamination. Huang et al. [Huang et al., 2025] construct hand authored hard perturbations of competition mathematics and attribute accuracy loss under an unchanged solution method to blind application of a learned procedure, a reading broadly consistent with the procedure level memorization in Section 4.2, established here against a corpus verified exactly rather than inferred. Wu et al. [Wu et al., 2025] show that reinforcement learning gains on mathematical reasoning depend on the pretraining exposure of the base policy and argue that clean synthetic problems are needed for reliable conclusions, a concern the checkpoint based design addresses by evaluating a policy before any relevant exposure occurred. Zhang et al. [Zhang et al., 2024] rebuild GSM8K from a contamination free procedure and report accuracy gaps of up to 13 points that track prior exposure across closed models, using an estimated exposure signal in the absence of any released training corpus; the present work substitutes an exact containment measurement for that estimate on a family where the corpus is public.

B Containment method, calibration, and corpus routes

The within split contamination contrast, high against low containment GSM8K train items, was closed rather than merely complicated by two findings. The corpus subset actually consumed during Instella 3B stage two pretraining, `train_119K`, is a 119,014 row subset of a 1,367,882 row released pool; containment computed against the pool over counts exposure, since 12.57 percent of GSM8K train reaches 0.80 or higher against the subset actually consumed against 23.34 percent against the pool, both measured on the same 891 parent problems through one extraction path. No GSM8K test item reaches containment 0.50 against either corpus, and the maximum over the full pool is 0.487, so the released corpus is train derived and the split remains a usable exposure contrast. Independently, `allenai/tulu-3-sft-mixture` is itself part of the Instella stage two pretraining mixture and `nvidia/OpenMathInstruct-2` enters at supervised fine tuning; both are generated in part from GSM8K train, and both carry it at containment 0.999 or higher on more than 97 percent of items. Every GSM8K train item therefore entered post stage one training through at least one of these routes regardless of its containment in the Instella specific synthetic corpus, so a high against low contrast defined on that corpus alone would compare two groups that are both almost entirely exposed once these mixtures are accounted for; the contrast is not attenuated by this, it is unidentifiable.

Item text is normalized to lowercase alphanumeric tokens and decomposed into every window of 13 consecutive tokens, following the overlap convention of Brown et al. [Brown et al., 2020]. Containment against a document is the fraction of an item’s windows found in that document; containment against a corpus is the maximum of that fraction over every document, computed in one streaming pass with an inverted index from window to source item so cost scales linearly in

corpus size. Because the corpus is generated only from the GSM8K training split, every GSM8K test item is a verified negative, and the false positive rate is measured rather than assumed: at the adopted threshold of 0.80, zero of 1319 test items are flagged, a Wilson 95 percent upper bound of 0.29 percent; a looser rule flagging any window match flags 3 of 1000 sampled test items, 0.30 percent.

Within the Tulu 3 mixture, a sampled prefix of 300,000 rows, six of the mixture’s subsets, finds GSM8K train items reaching containment 0.999 through the `flan_v2_converted` subset, 6021 of 7473 items from 89,982 rows of that source, and through `wildchat` conversation logs, 4 items from 100,000 rows, an organic route in which users pasted benchmark problems into chat. Every scan reports rows examined alongside items found, since a sampled prefix bounds containment from below and a zero means absent from the rows scanned rather than absent from the corpus.

C The irrelevance marker

The inserted frame in the crossed design ends "and has no bearing on this question", which is an instruction to ignore the insertion carried inside the insertion itself. Shi et al. [Shi et al., 2023] report that such an instruction mitigates the loss irrelevant context causes, so the modest cost measured in Section 3.3 may sit inside their mitigated regime rather than in tension with it. Testing that requires removing the clause and nothing else. The replacement, "on the afternoon of the same day", holds the sentence at exactly twenty words and preserves position, topic noun and unused integer.

Table 2: Accuracy against each problem’s own original, and the effect of removing the irrelevance marker. Cluster bootstrap over parent problems, 4000 draws, 1764 parents.

Checkpoint	Marked	Unmarked	Marker effect
Instella 3B	−2.21 [−3.88, −0.51]	−5.67 [−7.43, −3.77]	−3.46 [−4.90, −2.01]
SFT	+0.09 [−1.53, +1.67]	+0.11 [−1.53, +1.81]	+0.03 [−1.33, +1.36]
Instruct	−2.21 [−3.71, −0.65]	−2.78 [−4.34, −1.10]	−0.57 [−1.93, +0.79]

The marker governs the cost at the base stage two checkpoint and not after post training. One checkpoint of three is a single result rather than a pattern, and it is reported as such: it accounts for the size of the stage two effect and does not account for the residual two points at Instruct, which neither the number, nor topicality, nor the marker explains.

D The two withdrawn checks, and the dose response in full

The clause ladder is an externally authored perturbation suite whose rungs share template identifiers, which invites pairing each perturbed item against its own template head. The instances at each rung are independently generated rather than derived from one another, so a later rung is a different and harder problem set rather than the same problems with clauses appended, and an accuracy difference across rungs therefore mixes added content with added difficulty. Premise reordering permutes the non question sentences of a problem and preserves the gold answer by construction, but a narrative word problem often carries pronouns, temporal connectives and back references across sentences, so a permutation can leave text that no reader could solve. Filtering to items where no premise after the first carries such a dependency retains 24 percent of the set, and on that subset the effect is +2.2 points on [−6.5, +10.9] over 46 problems. Manual inspection finds violations that survive the filter, so even that residual is a spot check rather than an estimate.

Table 3 gives the dose response in full. Two outcomes are scored on the same generations. Answer recall is reproduction of the deleted parent’s original answer; procedure reproduction is the fraction of the parent’s intermediate calculator results appearing in the generation, excluding the final answer. Both are differences in differences, injected minus held out at a dose, minus the same quantity at dose zero, so the 8,388,608 tokens every arm trains on cancel. The dose zero row is the reference rather than a result, and its positive controls fail by construction, since a model that was never shown the documents cannot have memorised them.

Answer recall is flat in every model at every dose, including the two arms driven to near perfect verbatim reproduction, which is the result the body reports. Procedure reproduction rises with dose on Instella and reaches an interval excluding zero at the top dose only. On Qwen2.5 1.5B and OLMo

Table 3: Continued pretraining at five doses on four base models. Verbatim is the fraction of injected documents the model continues exactly after training. Answer and procedure are differences in differences against the dose zero arm, in percentage points, with cluster bootstrap intervals over parent problems, 4000 draws, seed 6198. Only one cell excludes zero, marked *. Qwen2.5-3B was run at two doses and is not scored for procedure reproduction.

Model	Dose	Verbatim	Answer recall	Procedure reproduction
Instella 3B	0×	0.133	reference	reference
	1×	0.153	+0.06 [−2.99, +3.17]	+0.34 [−2.86, +3.67]
	4×	0.180	+0.09 [−2.48, +2.68]	+0.73 [−2.27, +3.90]
	16×	0.491	−0.22 [−3.56, +3.10]	+2.78 [−0.79, +6.42]
	64×	0.946	+1.52 [−2.12, +5.21]	+5.44 [+1.46, +9.38]*
Qwen2.5 3B	0×	0.277	reference	not scored
	16×	0.713	+1.41 [−1.71, +4.60]	not scored
	64×	1.000	+2.28 [−1.21, +5.77]	not scored
Qwen2.5 1.5B	0×	0.243	reference	reference
	1×	0.250	+0.28 [−1.99, +2.60]	−2.17 [−5.05, +0.65]
	4×	0.255	+1.95 [−0.85, +4.84]	−0.70 [−3.69, +2.32]
	16×	0.406	+0.84 [−2.13, +3.71]	+0.83 [−2.65, +4.38]
	64×	0.989	+0.00 [−2.46, +2.52]	+0.33 [−3.49, +4.16]
OLMo 2 1B	0×	0.126	reference	reference
	1×	0.128	−0.32 [−2.38, +1.67]	−1.40 [−3.73, +0.93]
	4×	0.172	+0.00 [−1.90, +1.77]	−0.52 [−3.32, +2.30]
	16×	0.207	+0.62 [−1.56, +2.83]	+0.26 [−3.02, +3.49]
	64×	0.727	+0.28 [−2.34, +2.73]	−0.30 [−4.29, +3.61]

2 1B it does not, and the point estimates there are not consistently positive, so the chain result should be read as established on one model rather than replicated. The two replication models also memorise less at the same dose, reaching verbatim 0.989 and 0.727 against Instella’s 0.946 at 64× but only 0.406 and 0.207 at 16×, so the doses at which the Instella effect appears are not matched in memorisation strength. Distinguishing a genuine family difference from insufficient memorisation would need doses set to match verbatim reproduction rather than repetition count, which is not what these runs hold fixed.

E Robustness of the licence to its wording

The result in Section 3.1 is produced by one sentence, so it is worth knowing how much of it belongs to that sentence. The licence was reworded to carry the same three instructions, work stepwise, mark the final answer with the same token, and decline with ##### unanswerable when the problem does not fix a unique answer, in different words. The abstention token is held fixed so the scorer is unchanged between conditions and any movement is attributable to wording rather than to what is counted. All 3179 pruned items and all 1779 answerable parents were regenerated on all four checkpoints.

Table 4: The developmental ladder under two wordings of the same licence. Discrimination is declining on underdetermined items minus declining on answerable ones, in percentage points; false decline is the rate on answerable parents. Cluster bootstrap over parent problems, 4000 draws, seed 6198. Every interval excludes zero in both conditions.

Checkpoint	Discrimination		Decision quality		False decline	
	Original	Reworded	Original	Reworded	Original	Reworded
Stage 1	+1.84	+0.80	0.509	0.504	0.025	0.007
Instella 3B	+10.10	+6.64	0.550	0.533	0.008	0.008
SFT	+38.08	+33.41	0.690	0.667	0.035	0.008
Instruct	+52.63	+43.77	0.763	0.719	0.011	0.006

Two things separate. The developmental claim is robust to the wording: the ladder is monotone under both licences, in the same order, with every interval excluding zero, and false declines stay below one percent throughout the reworded condition. The level is not. Discrimination at instruct falls from +52.63 to +43.77 points and declining on underdetermined items falls from 0.538 to 0.444, so the rates quoted in Section 3.1 should be read as specific to the licence that produced them rather than

as properties of the checkpoint. The gap between conditions widens as the checkpoint improves, running -1.03 , -3.46 , -4.66 and -8.87 points, so the checkpoints that verify best are also the most sensitive to how they are asked. Truncation does not explain the shift, being equal or lower under the reworded licence at every checkpoint, and markedly lower at stage one, at 0.133 against 0.211 on pruned items.

The same rewording was applied to the cross family sweep of Section 3.3, on the identical items (Table 5). All four instruct checkpoints still discriminate with every interval excluding zero, and the ordering by decision quality is unchanged, so the comparison in Section 3.3 does not depend on the licence used to obtain it. Levels fall under rewording here as well. The false decline rate is the cell that moves most and it does not move uniformly: Qwen2.5 3B rises from 0.043 to 0.188 and OLMo 2 falls from 0.292 to 0.191 , while Instella stays near zero at 0.006 . The spread in how models respond to the wording of a licence is therefore wider than the spread in how well they discriminate, which is a further reason to report false refusal separately rather than folding it into a single rate.

Table 5: The cross family sweep under both wordings, identical GSM8K items. Every interval excludes zero in both conditions and the ordering by decision quality is preserved.

Instruct checkpoint	Discrimination		Decision quality		False decline	
	Original	Reworded	Original	Reworded	Original	Reworded
Qwen2.5 3B	+58.92	+45.33	0.795	0.727	0.043	0.188
Instella 3B	+52.63	+43.77	0.763	0.719	0.011	0.006
Qwen2.5 1.5B	+51.88	+39.83	0.759	0.699	0.090	0.061
OLMo 2 1B	+14.68	+9.17	0.573	0.546	0.292	0.191

F Reproduction

All generation, containment, and analysis code, together with every scored generation file and containment record underlying this work, is released for audit at a public repository and a paired dataset mirror, both anonymized for review and to be de-anonymized on acceptance. Every containment number reported uses the maximum over individual documents definition stated above; an earlier internal computation of the same quantity against the corpus pooled as one text produced a materially different figure for the same released split and is superseded by the document level definition throughout this work. Random seeds are fixed across item selection, deletion variant construction, the continued pretraining mixture, and every cluster bootstrap. An assistant was used to help implement generation, containment, and analysis code under the direction of the authors, who independently verified every reported number against the underlying generation files; the assistant did not design the probes or select which results to report.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state exactly the results reported in Sections 3 and 4: the developmental ladder in premise verification, its decay with reasoning depth, the replications on a second domain and three further families, and the dissociation between procedure and answer under controlled injection. Scope is stated as one model family at one parameter scale, and the injection result as one problem generator.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The Limitations paragraph of Section 5 states the scope limits explicitly: one problem generator for the injection arm, detection measured only through an explicit abstention option, difficulty stratified by the model’s own accuracy rather than an independent measure, depth on MATH measured as sentence count, chain reproduction under injection established on one model with the replication underpowered for it, and no attribution to specific training documents.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best

judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports empirical measurements and does not present theoretical results or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 2 specifies the corpus, containment method, checkpoint set, deletion probe construction, injection budget and doses, and perturbation constructions in enough detail to reproduce every reported number.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.

- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All generation, containment, and analysis code together with every scored generation file is released at an anonymized repository and dataset mirror referenced in Appendix D.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 2 specifies the containment threshold, deletion guard conditions, injection token budget and doses, decoding temperature, and item counts used for every reported result.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every reported effect in the tables and in Sections 3 and 4 is accompanied by a 95 percent cluster bootstrap interval over parent problems, 4000 draws at seed 6198, with the resampling unit stated in each table caption because deletion variants descending from one problem are not independent.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix D and Section 2 describe the generation engine, decoding settings, and the fixed token budget used for the continued pretraining intervention; exact wall clock and device counts are omitted and are available in the released repository logs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research measures properties of a publicly released language model family on a public arithmetic benchmark and introduces no new model, no human subjects, and no data collection beyond public benchmark items.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Broader impact is limited: the work concerns evaluation methodology for benchmark contamination and does not release a new capability; this is noted briefly in the conclusion.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No new model or dataset with misuse potential is released; the work analyzes existing public checkpoints and a public benchmark.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: Instella is credited and cited as Liu et al. 2025 with its released license noted in Appendix D; GSM8K, OpenMathInstruct-2, and the Tulu 3 mixture are each cited at first use in Sections 2 and 4 with their original authors.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The premise deletion item sets, the perturbation variant sets, and the injection mixture are new assets, documented in Section 2 and released with generation code in the accompanying repository.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The research involves no crowdsourcing and no human subjects; all evaluated text is machine generated against public benchmark items.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human subjects research was conducted.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: Appendix E discloses that an assistant was used to help implement generation, containment, and analysis code under the direction of the authors, who independently verified every reported number against the underlying generation files.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.