

---

# Language Models Reproduce Reasoning Procedures but Not Answers Under Verified Contamination

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

A benchmark item found in a model’s training data is routinely taken to mean its score on that item reflects recall, but testing the inference requires knowing what the model read, which the corpora of deployed systems do not permit. Instella publishes its full pretraining mixture and checkpoints from every stage, and its stage two corpus derives from the GSM8K training split, so membership is fixed by exact thirteen gram containment rather than estimated: against the 1319 test items, known negatives because the corpus derives only from train, the adopted threshold flags none, a Wilson upper bound of 0.29 percent. A premise deletion probe over 1591 train and 1588 test variants, across 891 and 888 parent problems, finds no rise in answer reproduction where the corpus enters training: a difference in differences of  $-0.26$  points on  $[-1.66, +1.14]$ , holding at  $+0.25$  and  $+0.94$  across the later boundaries, and unchanged under restriction to small answers, to terminated generations, and to a five bin containment gradient within the train arm that runs 3.7, 1.5, 0.9, 4.3 and 4.1 percent with no monotone order. Continued pretraining at five doses under a fixed token budget explains the null and separates two quantities usually conflated: at 64 repetitions the model reproduces 94.6 percent of injected text verbatim, yet against a never injected control the memorized solution procedure returns at  $+4.76$  points  $[+1.30, +8.23]$  while the final answer does not, at  $+1.52$  points  $[-1.93, +5.08]$ , procedure rising monotonically with dose and answer flat throughout. Exposure is legible in the calculation steps and not the answer, and needs far more repetition than the corpus supplies to any item, where document multiplicity reaches four. Two further checks locate what does cost accuracy: appending a dependent step to problems the strongest checkpoint already solves leaves 88.0 percent correct with zero reversions in 418 trials, while a crossed design holding sentence length, syntax and position fixed finds the main effect of an inserted unused number to be null at every checkpoint,  $-0.59$   $[-1.62, +0.45]$ ,  $-0.79$   $[-1.93, +0.31]$  and  $-0.42$   $[-1.53, +0.68]$ , with topicality null too and a residual cost of about two points for any insertion at all. Two perturbation designs indicating heavy fragility, an externally authored clause ladder and premise reordering, did not survive verification and are reported as withdrawn. Intervals are cluster bootstraps over parent problems, 4000 draws.

## 1 Introduction

Finding a benchmark item in a model’s training data is, by convention, grounds for discounting the score on that item. The step from the one to the other is rarely tested, since a test needs both certainty about what the model read and access to it at more than one point in training. Neither is available for frontier systems. Detection has grown up as a literature of proxies instead, from n gram overlap [Brown et al., 2020] to perplexity based membership inference [Carlini et al., 2023, Shi et al., 2024], and plain rephrasing is enough to defeat the overlap rules [Yang et al., 2023]. All

39 of these estimate membership. None establishes it, and an estimated treatment label cannot tell a  
40 real null apart from an uninformative one.

41 Instella [Liu et al., 2025] is the exception this work relies on. AMD publishes the complete pretrain-  
42 ing mixture together with checkpoints from every stage, including the stage two synthetic mathemat-  
43 ics corpus built from the GSM8K training split [Cobbe et al., 2021]. Most released models expose  
44 only final weights; Instella exposes the text consumed at each stage, so whether a given problem  
45 entered training, and when, becomes a matter of direct containment. Every GSM8K test item is  
46 a known negative with respect to that corpus, which turns detector error into a measured quantity:  
47 at the threshold adopted here, zero of 1319 test items are flagged, a Wilson upper bound of 0.29  
48 percent.

49 Knowing an item was available to be memorized says nothing about whether memorizing it inflated  
50 the score, so the probe removes one premise the solution genuinely uses. The original answer is then  
51 no longer derivable, and a model that remembered the parent problem can still emit it while a model  
52 working from the remaining text cannot. Instella ships four checkpoints along one trajectory, stage  
53 one before the synthetic corpus appears, stage two where it enters, an instruction tuned checkpoint  
54 and a preference optimized one, so the same items go to the model on both sides of that boundary.  
55 A memorization account predicts the rate climbs once the corpus arrives. It does not. To find what  
56 the probe could have detected, selected items were injected by continued pretraining at controlled  
57 doses, and there verbatim reproduction and answer recall come apart sharply.

## 58 2 Setup

59 **Corpus and checkpoints.** The stage two corpus, amd/Instella-GSM8K-synthetic, ships a  
60 generated pool and a 119 thousand item subset, train\_119K, documented as the split actually con-  
61 sumed in stage two pretraining. Containment is measured against that subset as a maximum over  
62 individual documents, so a claim of exposure points at one training document and not at the corpus  
63 treated as a single pooled text. Four checkpoints are probed: Stage 1, which precedes the corpus; In-  
64 stella 3B, where it enters; SFT; and Instruct. Generation runs on vLLM [Kwon et al., 2023], greedy  
65 for the primary estimates, with a temperature sweep as a check.

66 **Premise deletion.** One sentence the calculator annotated solution consumes is dropped from each  
67 GSM8K train and test parent problem. Two confounds correlated with containment are handled at  
68 construction. A necessity check confirms the deleted value appears nowhere else in the prompt, and  
69 variants whose original answer survives elsewhere are thrown out, since emitting it there is copying,  
70 and that leakage runs about twice as often on high containment items. What remains is 1591 train  
71 and 1588 test rows over 891 and 888 parents. Hit rate counts variants where the model emits the  
72 deleted parent’s original answer.

73 **Controlled injection.** Calibration comes from injecting 200 verified clean GSM8K test parents,  
74 alongside 200 held out entirely, into Instella 3B by continued pretraining. The token budget is fixed  
75 across five doses of zero, one, four, sixteen and sixty four repetitions, so a higher dose displaces  
76 generic filler and does not buy extra compute. Injected document loss and verbatim continuation sit  
77 inside every arm as positive controls, since a silently failed run and a genuine absence of memoriza-  
78 tion both give a flat curve. The same probe runs on both sets at every dose.

79 **Composition and perturbation.** A depth extension probe appends one dependent calculation step  
80 to problems the model already gets right, changing the gold answer, and asks whether the step  
81 is composed or the model falls back on the original problem. Answer preserving perturbations,  
82 paraphrase, an inserted off topic sentence, entity substitution, and a sentence carrying one unused  
83 number, are each scored against the same problem’s own original.

## 84 3 Results

### 85 3.1 Recall across the training trajectory

86 Figure 1 gives the hit rate on GSM8K train, which sits in the stage two corpus, and on GSM8K test,  
87 absent from every corpus scanned, together with the difference in differences at each boundary. A




Figure 1: (a) Premise deletion recall by checkpoint for items inside the stage two corpus and for verified clean items; the dotted line marks where the corpus enters. (b) Difference in differences between the arms at each training boundary. (c) Recall within the train arm alone, binned by containment against the consumed corpus, which uses no train against test contrast. (d) Verbatim reproduction of injected documents against dose, confirming the intervention took effect. (e) Reproduction of the memorized solution procedure and of the final answer, each against the never injected control. (f) Crossed design isolating the inserted number: four twenty word insertions differing only in an off topic against a problem derived noun and a filler word against an unused integer, each against its own untouched problem. Error bars are 95 percent cluster bootstraps over parent problems, 4000 draws.

88 memorization account puts the gap between stage one and stage two. Per checkpoint the train arm  
 89 runs 1.13, 2.83, 3.08 and 3.58 percent against 1.64, 3.59, 3.34 and 3.15 percent on the test arm, over  
 90 1591 and 1588 rows.

91 No gap opens at the boundary, and what movement there is runs the wrong way. The difference  
 92 in differences comes to  $-0.26$  points on an interval of  $[-1.66, +1.14]$ , and it stays flat at every  
 93 later boundary too,  $+0.25$  into SFT and  $+0.94$  into Instruct. Three variations leave it where it is:  
 94 restricting to one and two digit answers, restricting to rows where both checkpoints terminated, and  
 95 replacing the train against test contrast with a five bin containment gradient inside the train arm  
 96 alone. That last one matters most, since it drops the two arm comparison entirely, and the bins run  
 97 3.7, 1.5, 0.9, 4.3 and 4.1 percent from least to most contained, with no order to them. Over 891  
 98 parent clusters the design resolves 1.89 points at 80 percent power. Every gap observed falls under  
 99 that.

### 100 3.2 Procedure and answer under controlled injection

101 Injection worked: verbatim reproduction climbs from 13.3 percent with no injection to 94.6 percent  
 102 at the top dose. Figure 1 sets two outcomes against the never injected control at each dose, whether  
 103 the deleted premise value comes back and whether the memorized intermediate calculation steps  
 104 come back. Procedure reproduction runs  $+0.48$ ,  $+0.32$ ,  $+2.10$  and  $+4.76$  points across the four  
 105 nonzero doses, the last on  $[+1.30, +8.23]$ ; answer reproduction runs  $+0.06$ ,  $+0.09$ ,  $-0.22$  and  
 106  $+1.52$ , the last on  $[-1.93, +5.08]$ .

107 Answer recall never separates from the control, even where the model has all but memorized the  
 108 source text. Procedure reproduction climbs steadily and clears zero at the top dose. So the model  
 109 can reproduce a document almost word for word and still not recover the quantity that document  
 110 deletes, or the answer that quantity fixes. Exposure shows up in the calculation steps. It takes far

111 more of it than the released corpus ever supplies to one item, where document repetition tops out at  
112 four.

### 113 3.3 Composition and surface sensitivity

114 If remembering is not the constraint, two candidates remain: the model cannot chain an extra step,  
115 or any change to a problem’s surface throws it. Neither holds. Given a problem Instruct already  
116 solves and one dependent step appended, it stays correct on 88.0 percent, 368 of 418, and never falls  
117 back on the pre extension answer.

118 An earlier comparison suggested the inserted number was what mattered, since a sentence carrying  
119 an unused quantity cost 6.2 points where a plain distractor cost 0.6. Those conditions also differed  
120 in length and topicality, so a crossed design separates the three, holding length, syntax and position  
121 fixed across all four cells. The main effect of the number, fixed as the primary endpoint before the  
122 run, is null everywhere:  $-0.59$  on  $[-1.62, +0.45]$  at stage two,  $-0.79$  on  $[-1.93, +0.31]$  after fine  
123 tuning and  $-0.42$  on  $[-1.53, +0.68]$  at Instruct, over 1764 parents each. Topicality is null too. What  
124 survives is the cost of insertion itself, in panel (f) of Figure 1, where all four cells sit near  $-2$  points  
125 with overlapping intervals at stage two and at Instruct. The supervised fine tuned checkpoint is the  
126 exception and shows no insertion cost at all.

### 127 3.4 Recognising an unanswerable problem

128 Every deletion item is underdetermined by construction, so declining is correct, but the prompt used  
129 above demands a number and offers no alternative. Repeating the identical items under a prompt  
130 differing only in naming ##### unanswerable as an option changes the picture. Abstention runs  
131 7.5, 11.1, 42.4 and 53.2 percent across Stage 1, Instella 3B, SFT and Instruct, against roughly 0.1,  
132 0.0, 0.4 and 2.1 percent when the option is withheld, so the near zero figure measures compliance  
133 rather than competence and post training buys a steadily growing ability to recognise insufficiency.  
134 The strongest checkpoint still answers 44 percent of problems that provably have none, every de-  
135 tected abstention is an explicit refusal rather than an unparseable generation, and train and test are  
136 indistinguishable throughout.

### 137 3.5 Two withdrawn checks

138 Two designs looked at first like heavy evidence of fragility and neither held up, so both are reported  
139 as withdrawn rather than dropped. An externally authored clause ladder appeared to cost over 30  
140 points, but its later rungs are independently generated harder problems rather than the same items  
141 with clauses added, so part of the drop is difficulty; Ivanova [Ivanova, 2024] makes the general form  
142 of the point. Premise reordering appeared to cost over 12 points, but its generator permutes premises  
143 without checking cross sentence dependency, and restricted to items free of such dependency the  
144 effect lands on  $+2.2$  points over  $[-6.5, +10.9]$ . Appendix C gives both in full.

## 145 4 The unavailable contrast

146 The obvious design contrasts GSM8K train items at high containment against those at low, and it was  
147 tried first. Two findings closed it off, both in Appendix B. The subset the model actually consumed,  
148 train\_119K, puts 7.56 percent of train items at containment 0.80 or above against 29.75 percent  
149 computed on the full released pool. More decisively, allenai/tulu-3-sft-mixture [Lambert  
150 et al., 2024] sits in the stage two mixture and nvidia/OpenMathInstruct-2 [Toshniwal et al.,  
151 2024] arrives at fine tuning, each carrying GSM8K train at containment 0.999 or above on over 97  
152 percent of items against under 0.1 percent for test, one of the Tulu 3 subsets involved being labelled  
153 decontaminated. Every train item therefore reached the model after stage one by some route, so a  
154 high against low split on the Instella corpus alone compares two saturated groups. The checkpoint  
155 design never sorts items by containment and escapes this.

156 **Limitations and conclusion.** One family, one scale near three billion parameters, one benchmark  
157 domain, and a dissociation shown at five doses rather than ruled out at any exposure. Within those  
158 bounds verified exposure does not raise how often an answer returns, and does raise how often the  
159 procedure does.

## References

- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few shot learners. In *Advances in Neural Information Processing Systems*, 2020. arXiv:2005.14165.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *International Conference on Learning Representations*, 2023. arXiv:2202.07646.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Kaixuan Huang, Jiacheng Guo, Zihao Li, Xiang Ji, Jiawei Ge, Wenzhe Li, Yingqing Guo, Tianle Cai, Hui Yuan, Runzhe Wang, Yue Wu, Ming Yin, Shange Tang, Yujia Huang, Chi Jin, Xiaotong Chen, Chi Zhang, and Mengdi Wang. MATH perturb: Benchmarking LLMs mathematical reasoning abilities against hard perturbations. In *International Conference on Machine Learning*, 2025. arXiv:2502.06453.
- Dagny Rickard Ivanova. On some fixable limitations of understanding the limitations of mathematical reasoning in LLMs, 2024. Technical note.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Symposium on Operating Systems Principles*, 2023. arXiv:2309.06180.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tülu 3: Pushing frontiers in open language model post training. *arXiv preprint arXiv:2411.15124*, 2024.
- Jiang Liu, Jialian Wu, Xiaodong Yu, Yusheng Su, Prakamya Mishra, Gowtham Ramesh, Sudhanshu Ranjan, Chaitanya Manem, Ximeng Sun, Ze Wang, Pratik Prabhanjan Brahma, Zicheng Liu, and Emad Barsoum. Instella: Fully open language models with stellar performance. *arXiv preprint arXiv:2511.10628*, 2025.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. In *International Conference on Learning Representations*, 2024. arXiv:2310.16789.
- Shubham Toshniwal, Wei Du, Ivan Moshkov, Branislav Kisacanin, Alexan Ayrapetyan, and Igor Gitman. OpenMathInstruct-2: Accelerating AI for math with massive open source instruction data. *arXiv preprint arXiv:2410.01560*, 2024.
- Mingqi Wu, Zhihao Zhang, Qiang Dong, Zhiheng Xi, Jun Zhao, Senjie Jin, Xuanjing Fan, Yuan Zhou, Han Lv, Ming Zhang, Yuchen Fu, Qi Liu, Sheng Zhang, and Qi Zhang. Reasoning or memorization? unreliable results of reinforcement learning due to data contamination. *arXiv preprint arXiv:2507.10532*, 2025.
- Shuo Yang, Wei-Lin Chiang, Lianmin Zheng, Joseph E Gonzalez, and Ion Stoica. Rethinking benchmark and contamination for language models with rephrased samples. *arXiv preprint arXiv:2311.04850*, 2023.
- Hugh Zhang, Jeff Da, Dean Lee, Vaughn Robinson, Catherine Wu, Will Song, Tiffany Zhao, Pranav Raja, Dylan Slack, Qin Lyu, Sean Hendryx, Russell Kaplan, Michele Lunati, and Summer Yue. A careful examination of large language model performance on grade school arithmetic. *arXiv preprint arXiv:2405.00332*, 2024.

## A Related work

Huang et al. [Huang et al., 2025] construct hand authored hard perturbations of competition mathematics and attribute accuracy loss under an unchanged solution method to blind application of a learned procedure, a reading broadly consistent with the procedure level memorization in Section 3.2, established here against a corpus verified exactly rather than inferred. Wu et al. [Wu et al., 2025] show that reinforcement learning gains on mathematical reasoning depend on the pretraining exposure of the base policy and argue that clean synthetic problems are needed for reliable conclusions, a concern the checkpoint based design addresses by evaluating a policy before any relevant exposure occurred. Zhang et al. [Zhang et al., 2024] rebuild GSM8K from a contamination free procedure and report accuracy gaps of up to 13 points that track prior exposure across closed models, using an estimated exposure signal in the absence of any released training corpus; the present work substitutes an exact containment measurement for that estimate on a family where the corpus is public.

## B Containment method, calibration, and corpus routes

The within split contamination contrast, high against low containment GSM8K train items, was closed rather than merely complicated by two findings. The corpus subset actually consumed during Instella 3B stage two pretraining, `train_119K`, is a 119,014 row subset of a 1,367,882 row released pool; containment computed against the pool over counts exposure, since only 7.56 percent of GSM8K train reaches 0.80 or higher against the subset actually consumed against 29.75 percent against the pool. Independently, `allenai/tulu-3-sft-mixture` is itself part of the Instella stage two pretraining mixture and `nvidia/OpenMathInstruct-2` enters at supervised fine tuning; both are generated in part from GSM8K train, and both carry it at containment 0.999 or higher on more than 97 percent of items. Every GSM8K train item therefore entered post stage one training through at least one of these routes regardless of its containment in the Instella specific synthetic corpus, so a high against low contrast defined on that corpus alone would compare two groups that are both almost entirely exposed once these mixtures are accounted for; the contrast is not attenuated by this, it is unidentifiable.

Item text is normalized to lowercase alphanumeric tokens and decomposed into every window of 13 consecutive tokens, following the overlap convention of Brown et al. [Brown et al., 2020]. Containment against a document is the fraction of an item’s windows found in that document; containment against a corpus is the maximum of that fraction over every document, computed in one streaming pass with an inverted index from window to source item so cost scales linearly in corpus size. Because the corpus is generated only from the GSM8K training split, every GSM8K test item is a verified negative, and the false positive rate is measured rather than assumed: at the adopted threshold of 0.80, zero of 1319 test items are flagged, a Wilson 95 percent upper bound of 0.29 percent; a looser rule flagging any window match flags 3 of 1000 sampled test items, 0.30 percent.

Within the Tulu 3 mixture, a sampled prefix of 300,000 rows, six of the mixture’s subsets, finds GSM8K train items reaching containment 0.999 through the `flan_v2_converted` subset, 6021 of 7473 items from 89,982 rows of that source, and through `wildchat` conversation logs, 4 items from 100,000 rows, an organic route in which users pasted benchmark problems into chat. Every scan reports rows examined alongside items found, since a sampled prefix bounds containment from below and a zero means absent from the rows scanned rather than absent from the corpus.

## C The two withdrawn checks in full

The clause ladder is an externally authored perturbation suite whose rungs share template identifiers, which invites pairing each perturbed item against its own template head. The instances at each rung are independently generated rather than derived from one another, so a later rung is a different and harder problem set rather than the same problems with clauses appended, and an accuracy difference across rungs therefore mixes added content with added difficulty. Premise reordering permutes the non question sentences of a problem and preserves the gold answer by construction, but a narrative word problem often carries pronouns, temporal connectives and back references across sentences, so a permutation can leave text that no reader could solve. Filtering to items where no premise after the first carries such a dependency retains 24 percent of the set, and on that subset the effect is +2.2

260 points on  $[-6.5, +10.9]$  over 46 problems. Manual inspection finds violations that survive the filter,  
261 so even that residual is a spot check rather than an estimate.

## 262 **D Reproduction**

263 All generation, containment, and analysis code, together with every scored generation file and con-  
264 tainment record underlying this work, is released for audit at a public repository and a paired dataset  
265 mirror, both anonymized for review and to be de anonymized on acceptance. Every containment  
266 number reported uses the maximum over individual documents definition stated above; an earlier  
267 internal computation of the same quantity against the corpus pooled as one text produced a materi-  
268 ally different figure for the same released split and is superseded by the document level definition  
269 throughout this work. Random seeds are fixed across item selection, deletion variant construction,  
270 the continued pretraining mixture, and every cluster bootstrap. An assistant was used to help imple-  
271 ment generation, containment, and analysis code under the direction of the authors, who indepen-  
272 dently verified every reported number against the underlying generation files; the assistant did not  
273 design the probes or select which results to report.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state exactly the four findings reported in Section 3 and the corpus limitation reported in Section 4, with no broader claim of generalization beyond one model family and grade school arithmetic.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 5 states the scope limits explicitly: one model family, one parameter scale, one benchmark domain, the untested unused number versus sentence length confound, and the unrun licensing prompt condition.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports empirical measurements and does not present theoretical results or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 2 specifies the corpus, containment method, checkpoint set, deletion probe construction, injection budget and doses, and perturbation constructions in enough detail to reproduce every reported number.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All generation, containment, and analysis code together with every scored generation file is released at an anonymized repository and dataset mirror referenced in Appendix D.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 2 specifies the containment threshold, deletion guard conditions, injection token budget and doses, decoding temperature, and item counts used for every reported result.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every reported effect in Tables 1 through 3 and in Section 3 is accompanied by a 95 percent cluster bootstrap interval over parent problems, 4000 draws, stated in each table caption.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix D and Section 2 describe the generation engine, decoding settings, and the fixed token budget used for the continued pretraining intervention; exact wall clock and device counts are omitted and are available in the released repository logs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research measures properties of a publicly released language model family on a public arithmetic benchmark and introduces no new model, no human subjects, and no data collection beyond public benchmark items.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Broader impact is limited: the work concerns evaluation methodology for benchmark contamination and does not release a new capability; this is noted briefly in the conclusion.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No new model or dataset with misuse potential is released; the work analyzes existing public checkpoints and a public benchmark.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Instella is credited and cited as Liu et al. 2025 with its released license noted in Appendix D; GSM8K, OpenMathInstruct-2, and the Tulu 3 mixture are each cited at first use in Sections 2 and 4 with their original authors.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The premise deletion item sets, the perturbation variant sets, and the injection mixture are new assets, documented in Section 2 and released with generation code in the accompanying repository.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research involves no crowdsourcing and no human subjects; all evaluated text is machine generated against public benchmark items.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

587 **15. Institutional review board (IRB) approvals or equivalent for research with human**  
588 **subjects**

589 Question: Does the paper describe potential risks incurred by study participants, whether  
590 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)  
591 approvals (or an equivalent approval/review based on the requirements of your country or  
592 institution) were obtained?

593 Answer: [N/A]

594 Justification: No human subjects research was conducted.

595 Guidelines:

- 596 • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
597 with human subjects.
- 598 • Depending on the country in which research is conducted, IRB approval (or equiva-  
599 lent) may be required for any human subjects research. If you obtained IRB approval,  
600 you should clearly state this in the paper.
- 601 • We recognize that the procedures for this may vary significantly between institutions  
602 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
603 guidelines for their institution.
- 604 • For initial submissions, do not include any information that would break anonymity  
605 (if applicable), such as the institution conducting the review.

606 **16. Declaration of LLM usage**

607 Question: Does the paper describe the usage of LLMs if it is an important, original, or  
608 non-standard component of the core methods in this research? Note that if the LLM is used  
609 only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
610 scientific rigor, or originality of the research, declaration is not required.

611 Answer: [Yes]

612 Justification: Appendix E discloses that an assistant was used to help implement genera-  
613 tion, containment, and analysis code under the direction of the authors, who independently  
614 verified every reported number against the underlying generation files.

615 Guidelines:

- 616 • The answer [N/A] means that the core method development in this research does not  
617 involve LLMs as any important, original, or non-standard components.
- 618 • Please refer to our LLM policy in the NeurIPS handbook for what should or should  
619 not be described.